

bitcoin-timeseries-ml-engineering (Public)

p ...

1 Branch

0 Tags

Go to file

Go to file

Add file

Code

...

About



No description, website, or topics provided.

Readme

Activity

0 stars

0 watching

0 forks

Releases

No releases published

[Create a new release](#)

Packages

No packages published

[Publish your first package](#)

Languages

Python 100.0%

Suggested workflows

Based on your tech stack



SLSA Generic generator

[Configure](#)

Generate SLSA3 provenance for your existing release workflows



Django

[Configure](#)

Build and Test a Django Project



Python package

[Configure](#)

Create and test a Python package on multiple Python versions.

[More workflows](#)[Dismiss suggestions](#)

README



bitcoin-timeseries-ml-engineering

Bitcoin Time-Series ML Engineering (LSTM + Walk-Forward Validation) — portfolio-grade pipeline with private alpha redaction

Portfolio project for the IBM AI Engineering Certificate

This repo demonstrates end-to-end ML engineering for financial time-series (data processing → model training → evaluation → reproducible runs).

Important boundary: my proprietary trading datasets, QFL-DCA rules, and "alpha discoveries" are intentionally **not** published.

What this repo is

A production-style ML pipeline for Bitcoin forecasting experiments using:

- Time-series-safe splits (chronological + walk-forward evaluation)
- Leakage-resistant scaling (fit on train only)
- Model architectures (LSTM variants + optional attention/CNN)
- Evaluation (standard ML metrics + trading-aware metrics)

This is designed to be auditable, reproducible, and recruiter-readable.

What this repo is NOT

- Not a "copy-paste profitable strategy"
- Not a data dump of my historical trades
- Not a release of my full QFL-DCA alpha rulebook
- Not financial advice

If you want to reproduce results, you must use **your own datasets** (or the optional public sample described below).

Privacy / Alpha Redaction Policy (Explicit)

To protect years of research and private trading data, the following are excluded from this public repo:

- Private data/ (raw, processed, or labeled datasets)
- Private outputs/ (full experiment result tables, matched trades, correlations)
- Certain strategy-specific configurations / thresholds

- Full QFL-DCA optimization documents or proprietary rule sets

What is included instead:

- The pipeline and engineering rigor
- Config-driven training/evaluation
- A clear "Bring Your Own Data" interface
- Optional synthetic or public sample dataset support (recommended)

Repo Structure (public-safe)

```
.
├── src/
│   ├── data/          # feature engineering + processor (train-only scaling)
│   ├── models/        # LSTM architectures
│   ├── training/     # trainer + metrics + walk-forward
│   └── utils/         # inference utilities
├── tests/           # unit tests (pipeline sanity checks)
└── notebooks/       # exploration / colab (sanitized)

├── configs/         # PUBLIC configs (no private thresholds)
├── docs/            # methodology notes (portfolio narrative)
├── scripts/         # helper scripts (sanitized)
├── main.py          # CLI entry point (train/evaluate)
├── requirements.txt
└── README.md        # contributor/agent runbook (hardening + release flow)

└── AGENT.md          # contributor/agent runbook (hardening + release flow)
```



Engineering Rigor Highlights (what recruiters should notice)

- Reproducibility: seed control + config-driven runs (config.yaml)
(see seed handling in main.py)
- Leakage control: scaling is fit on train split only (zero-leakage design)
- Time-series validation: walk-forward cross-validation (no random k-fold)
- Time-series rigor: DataLoaders use shuffle=False to preserve temporal order
- Checkpointing: best model is saved and reloaded for evaluation
- Metrics: ML metrics + trading-aware metrics (Sharpe, drawdown, directional accuracy)

Quick Start (Bring Your Own Data)

1) Install

```
python -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt
```



2) Provide your BTC data

You provide a CSV at minimum containing OHLCV daily bars.

Expected location (default):

- data/raw/btc_ohlcv_daily.csv (not committed)

Optional:

- data/raw/btc_onchain.csv (not committed)

If you don't want to use private data, use the optional "public sample mode" once you add it (see below).

3) Train / Evaluate

```
python main.py --mode train --config config.yaml
python main.py --mode evaluate --config config.yaml
```



Outputs (local only, not committed):

- models/best_model.pt
- models/scalers.joblib
- outputs/metrics.json / outputs/metrics.txt
- outputs/test_predictions.csv

Public Sample Mode (Recommended)

To keep this repo reproducible without exposing private data:

- Add a **small public dataset** (or generate synthetic data) under:
 - `data/sample/` (committed)
- Keep real datasets in:
 - `data/raw/` (`gitignored`)

This keeps the project runnable for reviewers while protecting your edge.

Results (How to interpret)

This repo focuses on **engineering quality** and correct evaluation for noisy markets.

Important:

- Any "paper numbers" (e.g., $R^2=0.991$) are **literature-reported benchmarks**, not guaranteed reproduction.
- Your actual results depend heavily on data window, feature set, regime shifts, and cost assumptions.

Safety / Compliance

This is educational and for portfolio demonstration only. No financial advice. Use at your own risk.

License

Choose a permissive license for portfolio visibility (MIT/Apache-2.0). If you want to restrict use, choose a more protective license. (Decide before first push.)

Contact

If you're a recruiter/hiring manager: this repo is intended to show ML engineering skill in time-series forecasting + evaluation discipline.